

Schema Design

schema, n. — a representation of a plan or theory in the form of an outline or model.

Schemas

- **Relation Schema (i.e. table schema)**
 - Logical definition of the relation (table)
- **Database Schema**
 - Blueprint of the collection of relation schemas for an entire database

Data Modeling

- How do we represent real world relationships and properties in our program?
 - ...in a way that makes writing the program easy
 - ...while remaining flexible for future changes
 - ...oh, it also has to be fast (enough).

Designing a Schema

- **Analysis**
 - What does my program need to output?
 - What data will I need to produce that output?
- **Conceptual Design**
 - Conceptual entities and their relationships
- **Logical Design**
 - In a SQL database: What are my tables, attributes, and relationships?
 - In a program: What are my functions and data structures?
- **Physical Design**
 - JavaScript code, CREATE TABLE statements

Designing a Schema

What we'll focus on today

- **Analysis**

- What does my program need to output?
- What data will I need to produce that output?

- **Conceptual Design**

- Conceptual entities and their relationships

- **Logical Design**

- In a SQL database: What are my tables, attributes, and relationships?
- In a program: What are my functions and data structures?

- **Physical Design**

- JavaScript code, CREATE TABLE statements

Designing a Schema

- **Analysis**
 - What does my program need to output?
 - What data will I need to produce that output?
- **Conceptual Design**
 - Conceptual entities and their relationships
- **Logical Design**
 - In a SQL database: What are my tables, attributes, and relationships?
 - In a program: What are my functions and data structures?
- **Physical Design**
 - JavaScript code, CREATE TABLE statements

Example: A Journal Analysis


- I want a program to keep my journal in.
- I want to be able to enter the **text** of each journal entry.
- I want to be able to see journal entries **chronologically**.

Designing a Schema

- **Analysis**
 - What does my program need to output?
 - What data will I need to produce that output?
- **Conceptual Design**
 - Conceptual entities and their relationships
- **Logical Design**
 - In a SQL database: What are my tables, attributes, and relationships?
 - In a program: What are my functions and data structures?
- **Physical Design**
 - JavaScript code, CREATE TABLE statements

Entity Relationship Diagram (ERD)

Conceptual Design



Designing a Schema

- **Analysis**
 - What does my program need to output?
 - What data will I need to produce that output?
- **Conceptual Design**
 - Conceptual entities and their relationships
- **Logical Design**
 - In a SQL database: What are my tables, attributes, and relationships?
 - In a program: What are my functions and data structures?
- **Physical Design**
 - JavaScript code, CREATE TABLE statements

Entity Relationship Diagram (ERD)

Logical Design

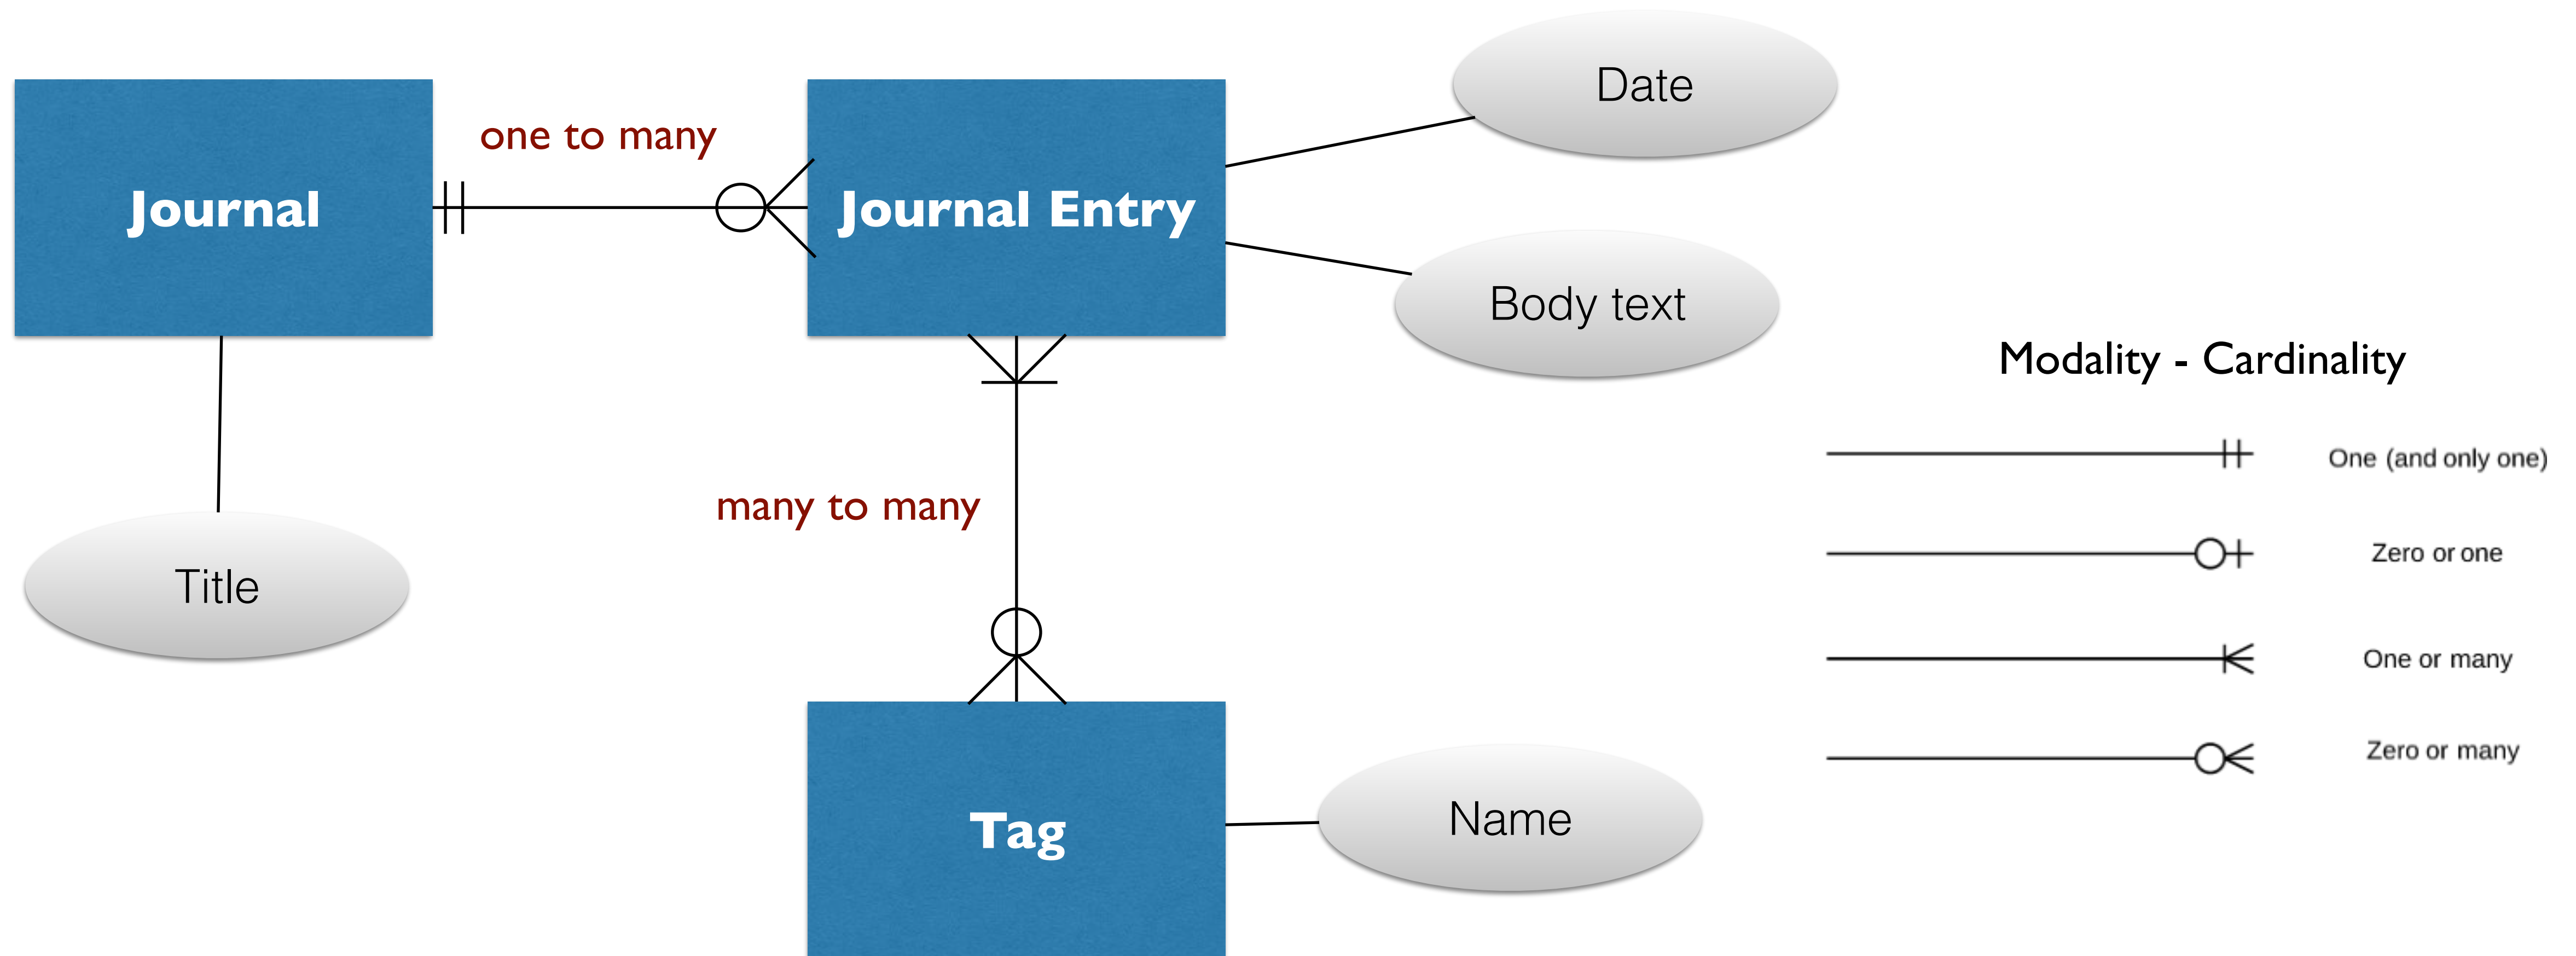
entries	
id	int, primary key
date_created	date
text	text

All done!

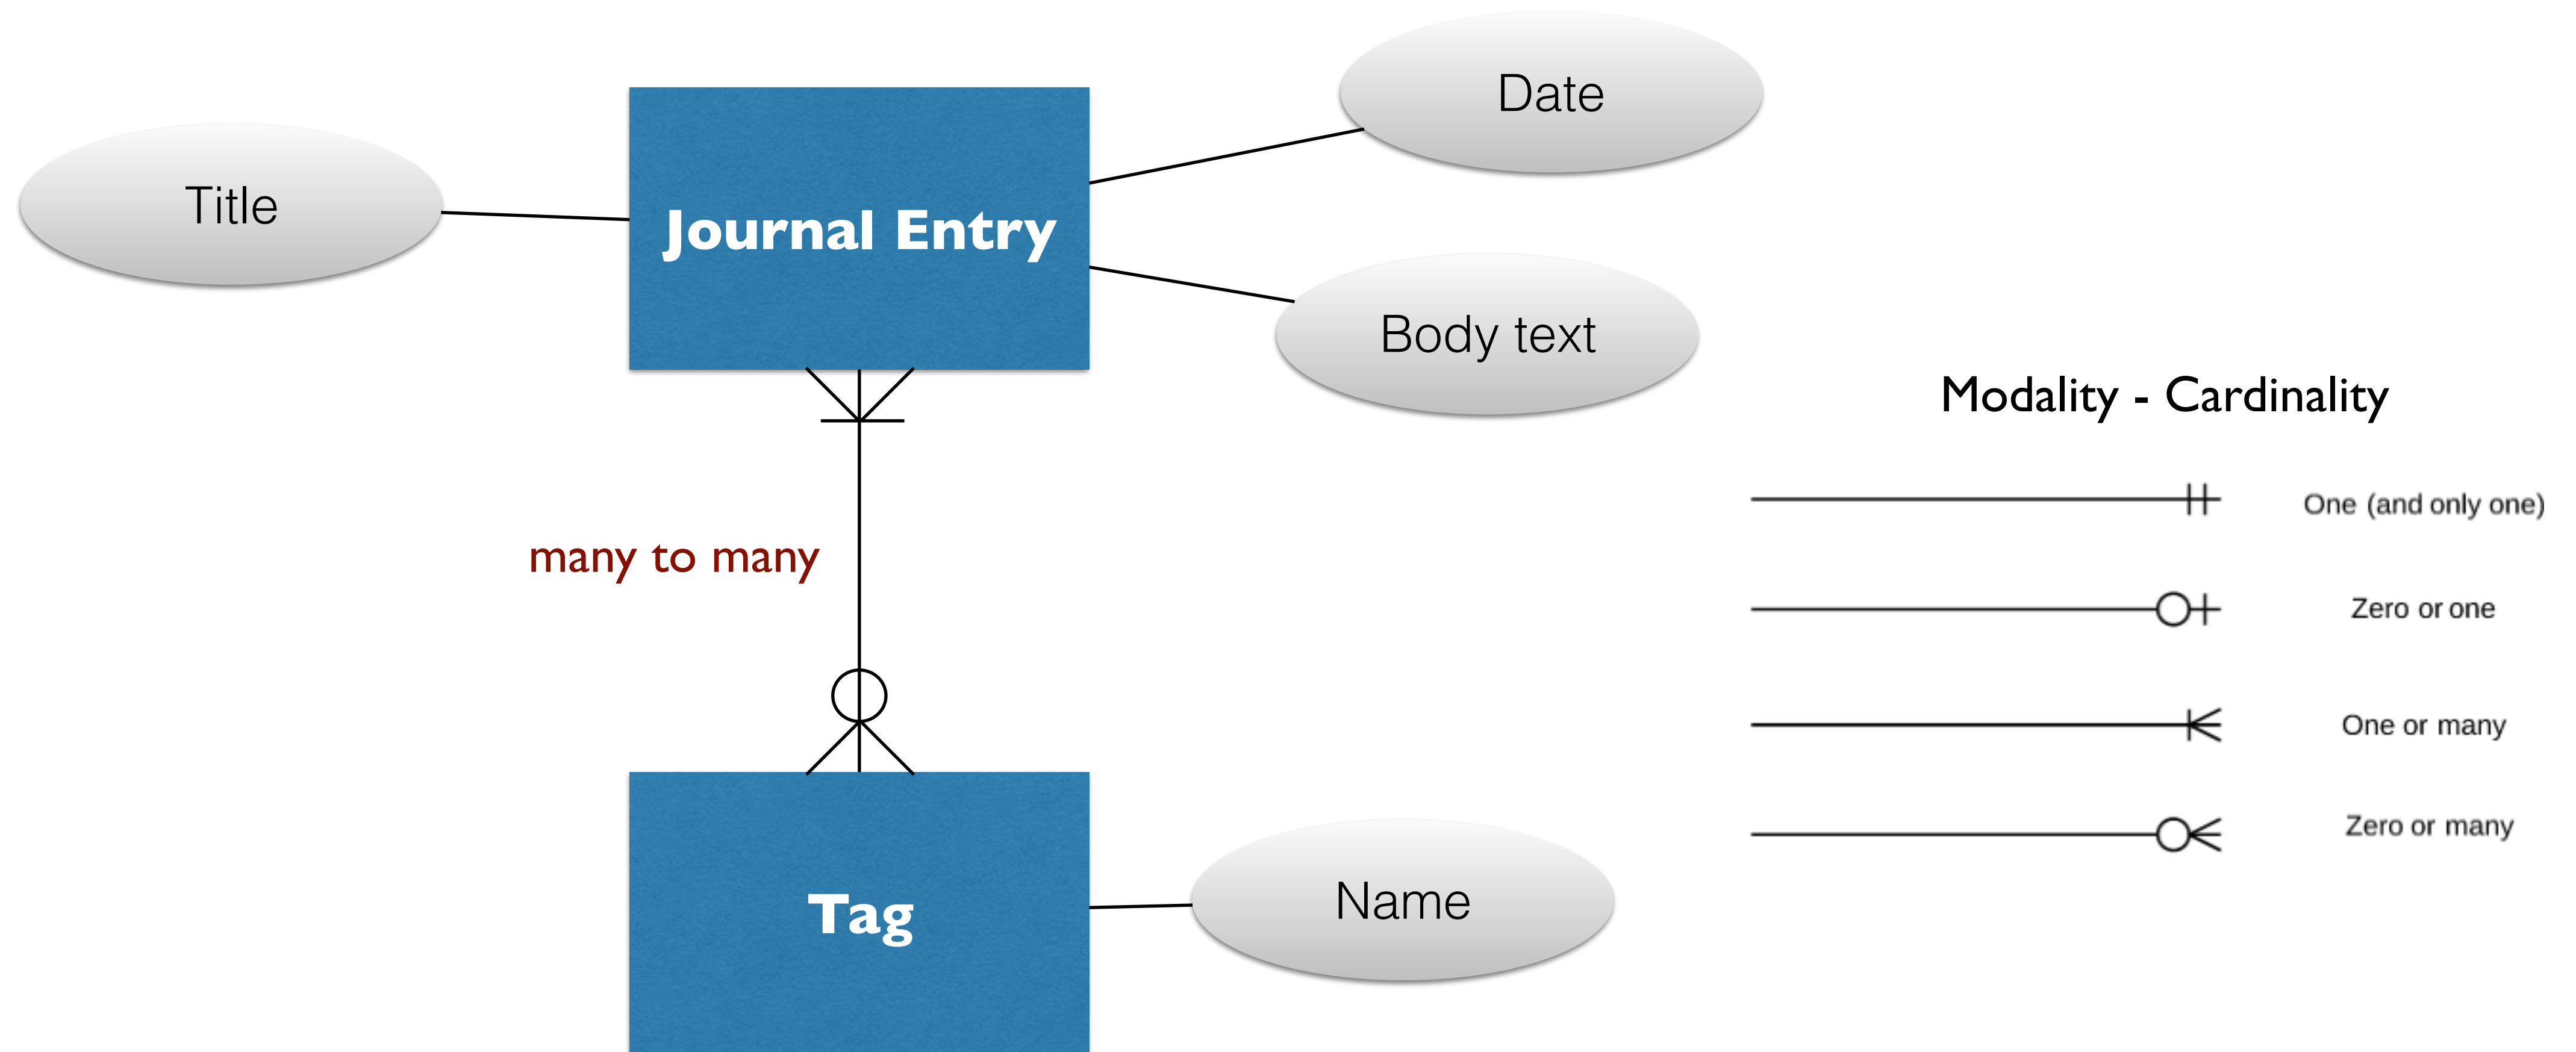
- Oh wait, I forgot a couple of things
 - I want to be able to have multiple journals
 - I want to be able to #tag entries and find all entries with a particular #tag
- Take 2...

Example: A Journal

Conceptual Design, Take 2



...Or maybe like this?

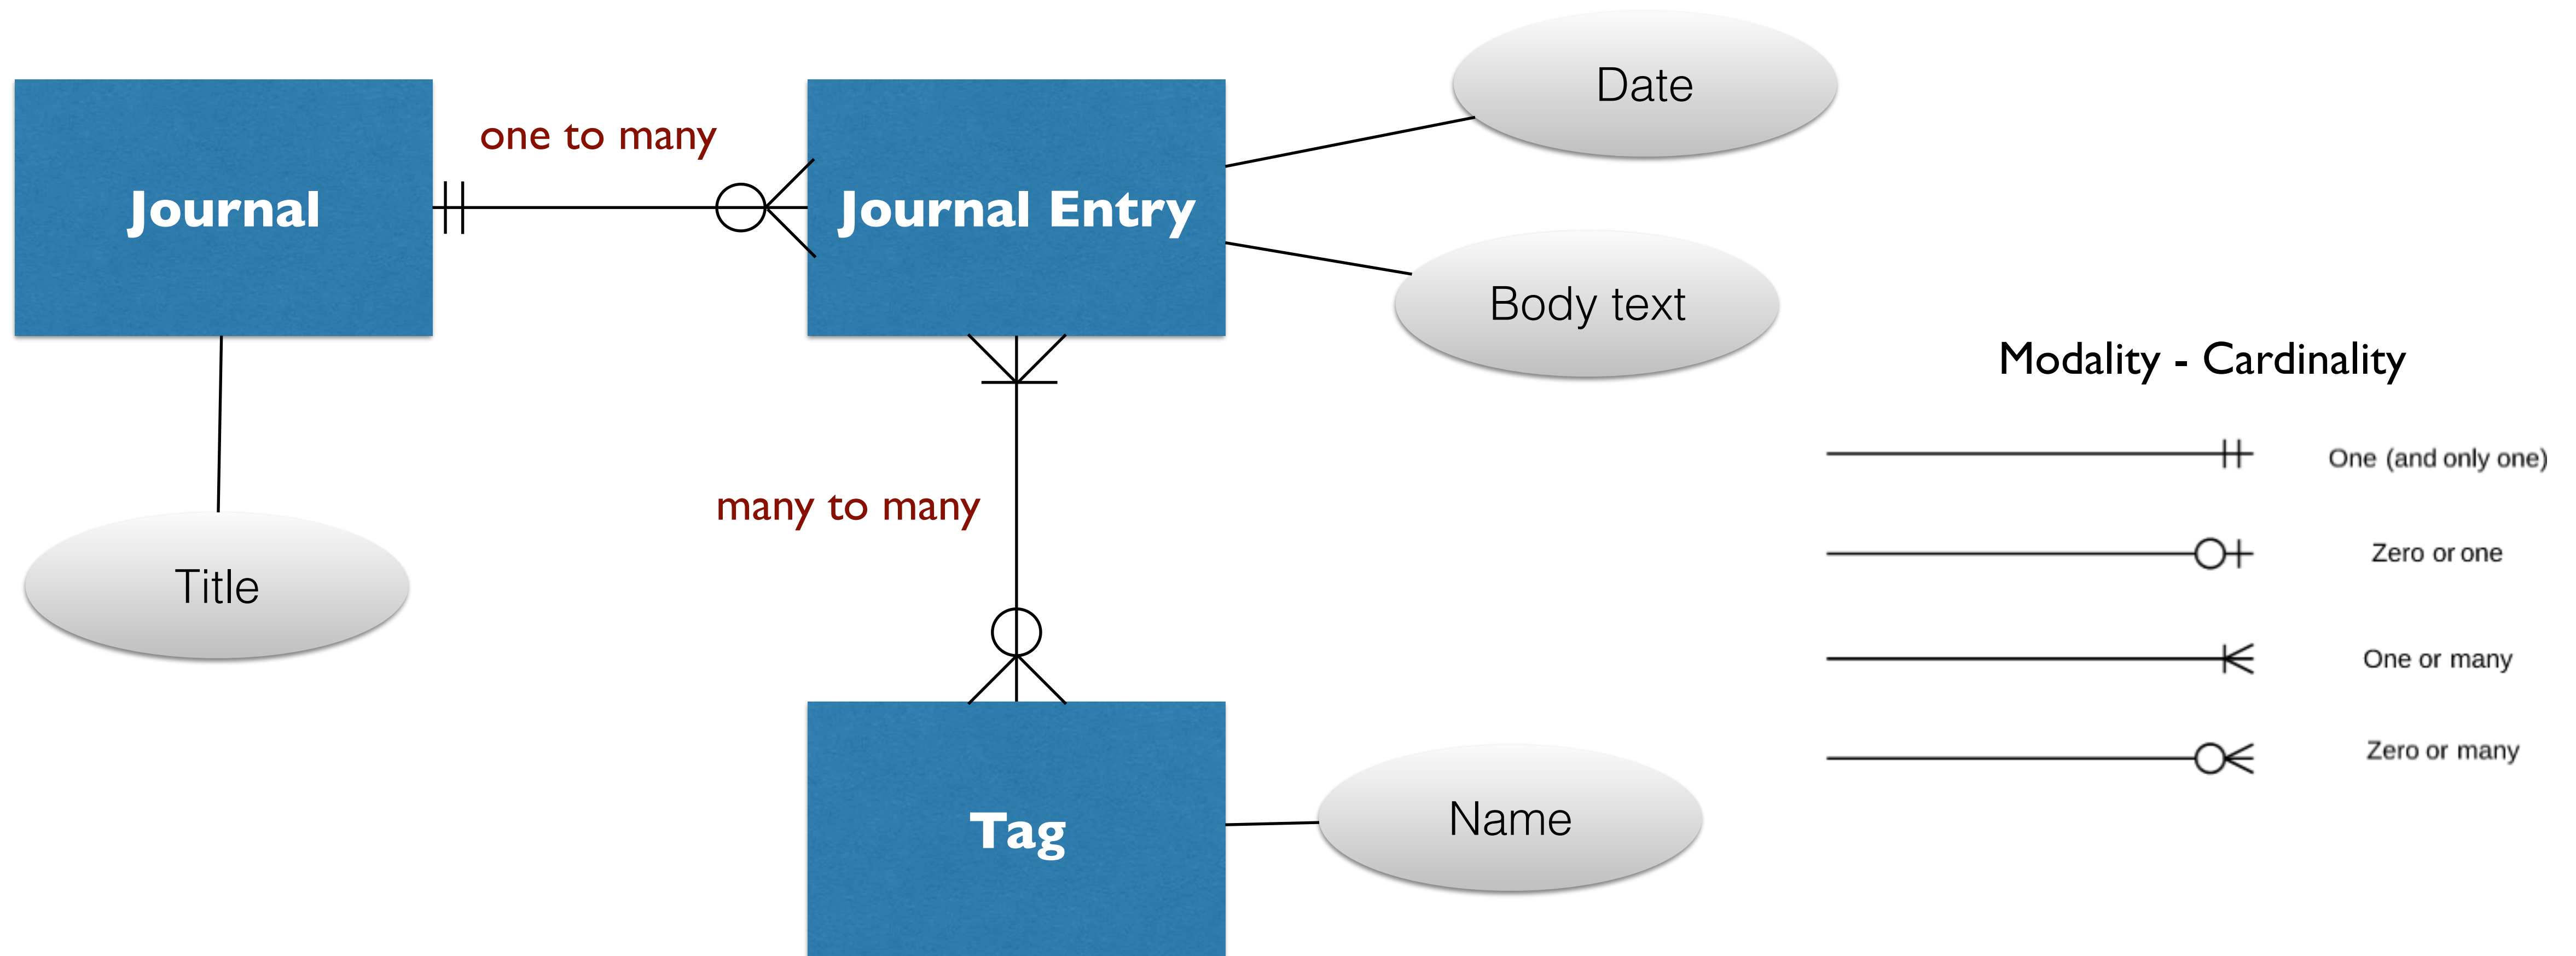


- How do I change the name of “happy times” to “sadness”?

select * from entries;			
id	date_created	text	journal_title
0	2016-04-01	I am happy	happy times
1	2016-04-02	I am very happy	happy times
2	2016-04-03	Despair fills me	happy times
3	2016-04-03	Sadness is my life	an anatomy of pain


Example: A Journal

Conceptual Design, Take 2




Example: A Journal

Logical Design: Take 2?



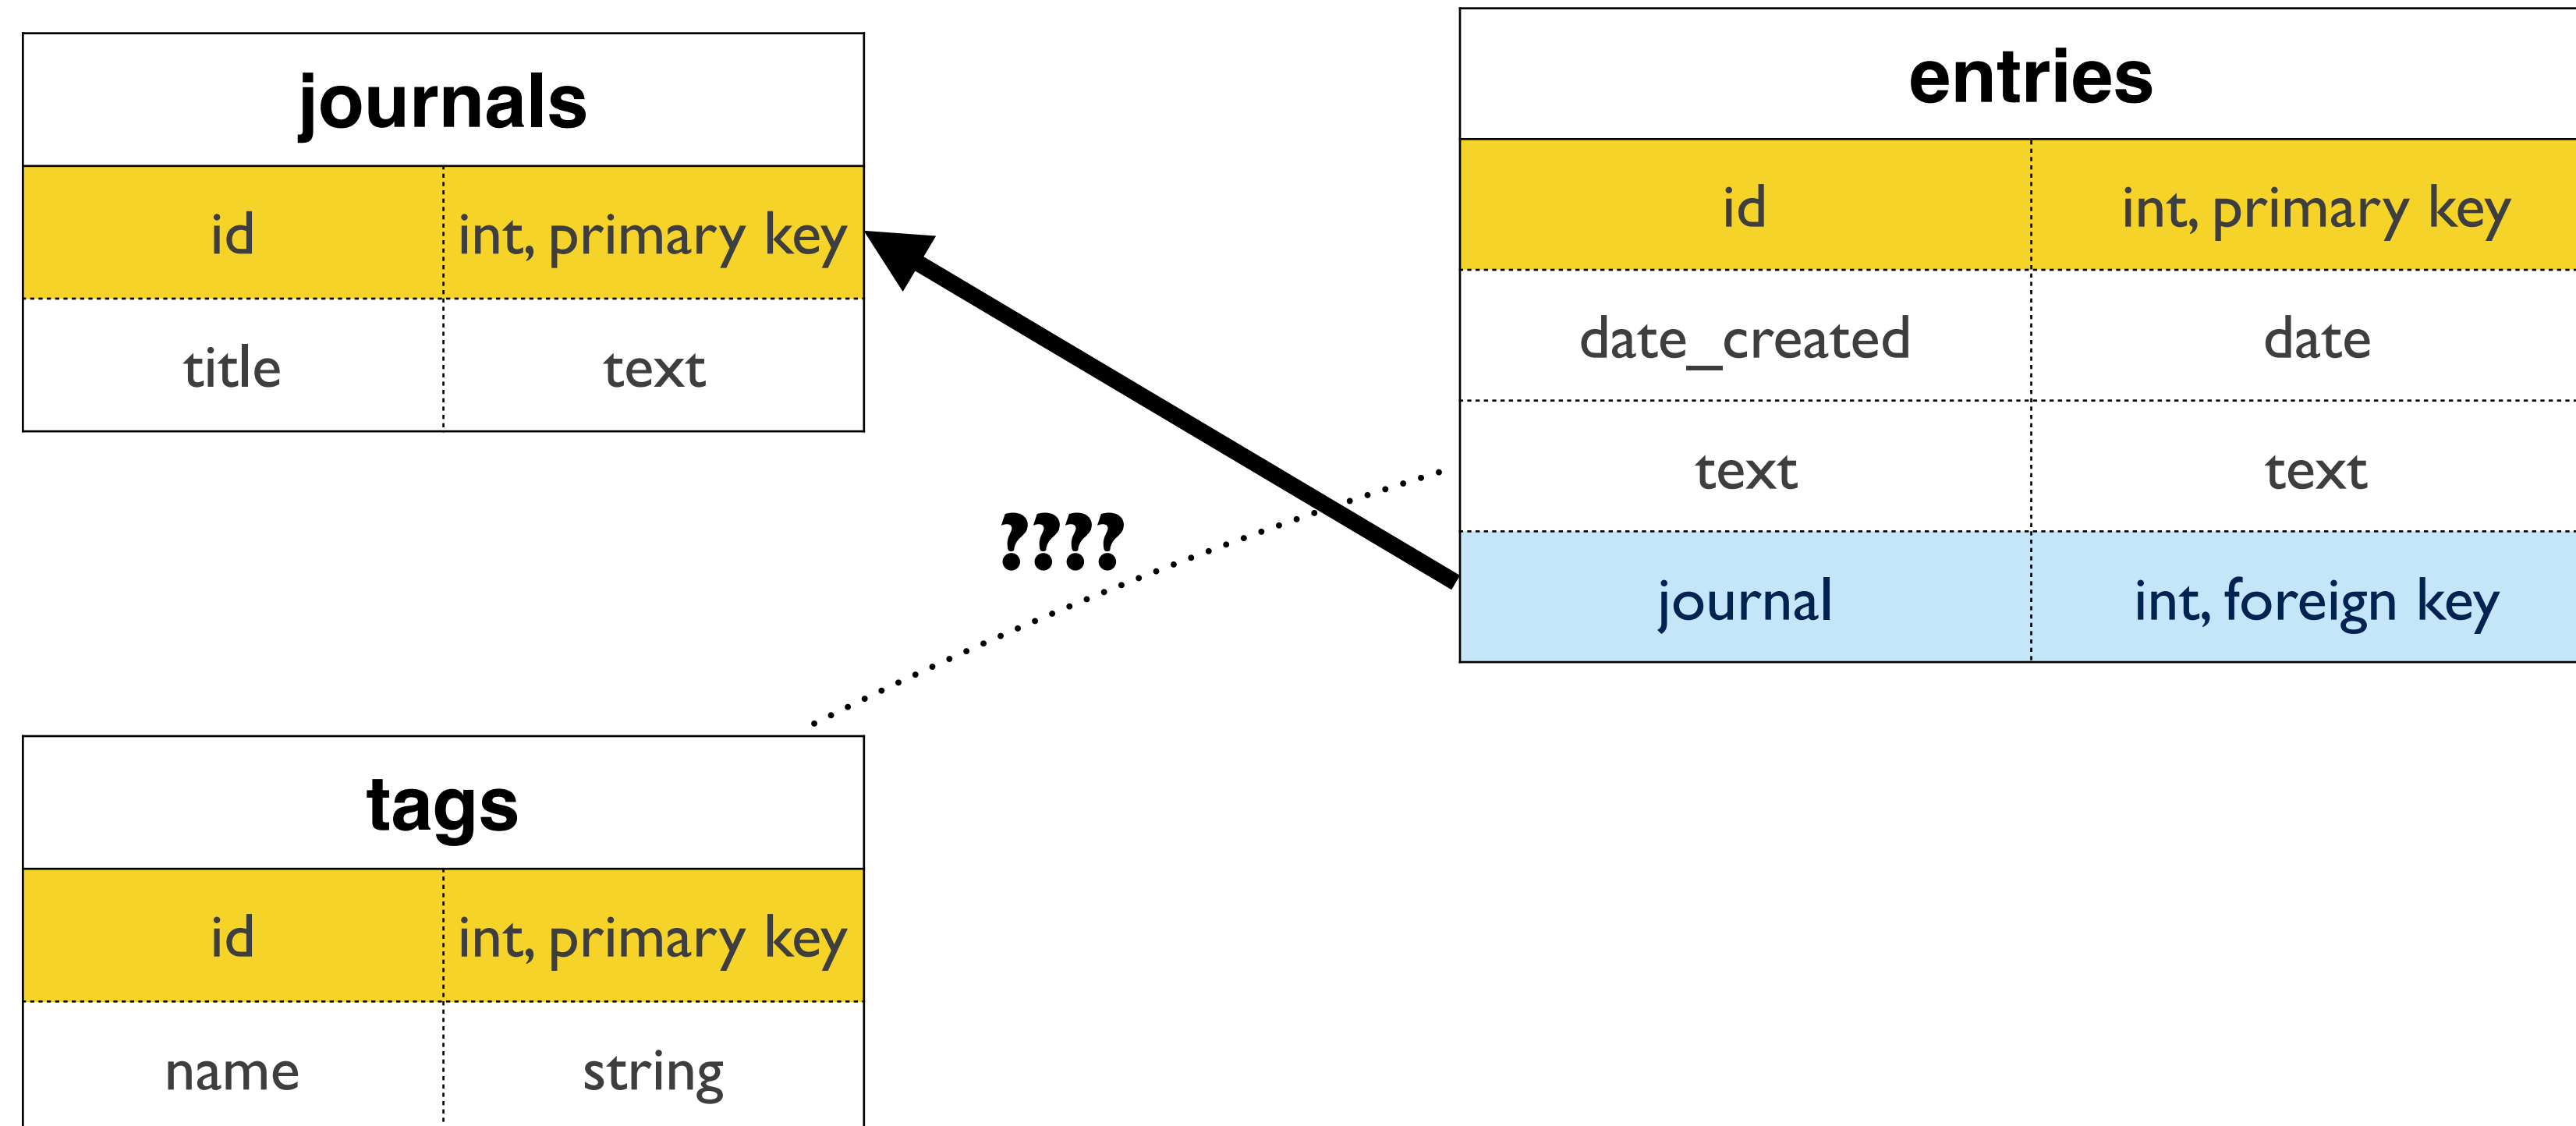
Normalization

- Eliminate repeating groups in individual tables
- Create a separate table for each set of related data
- Identify each set of related data with a primary key



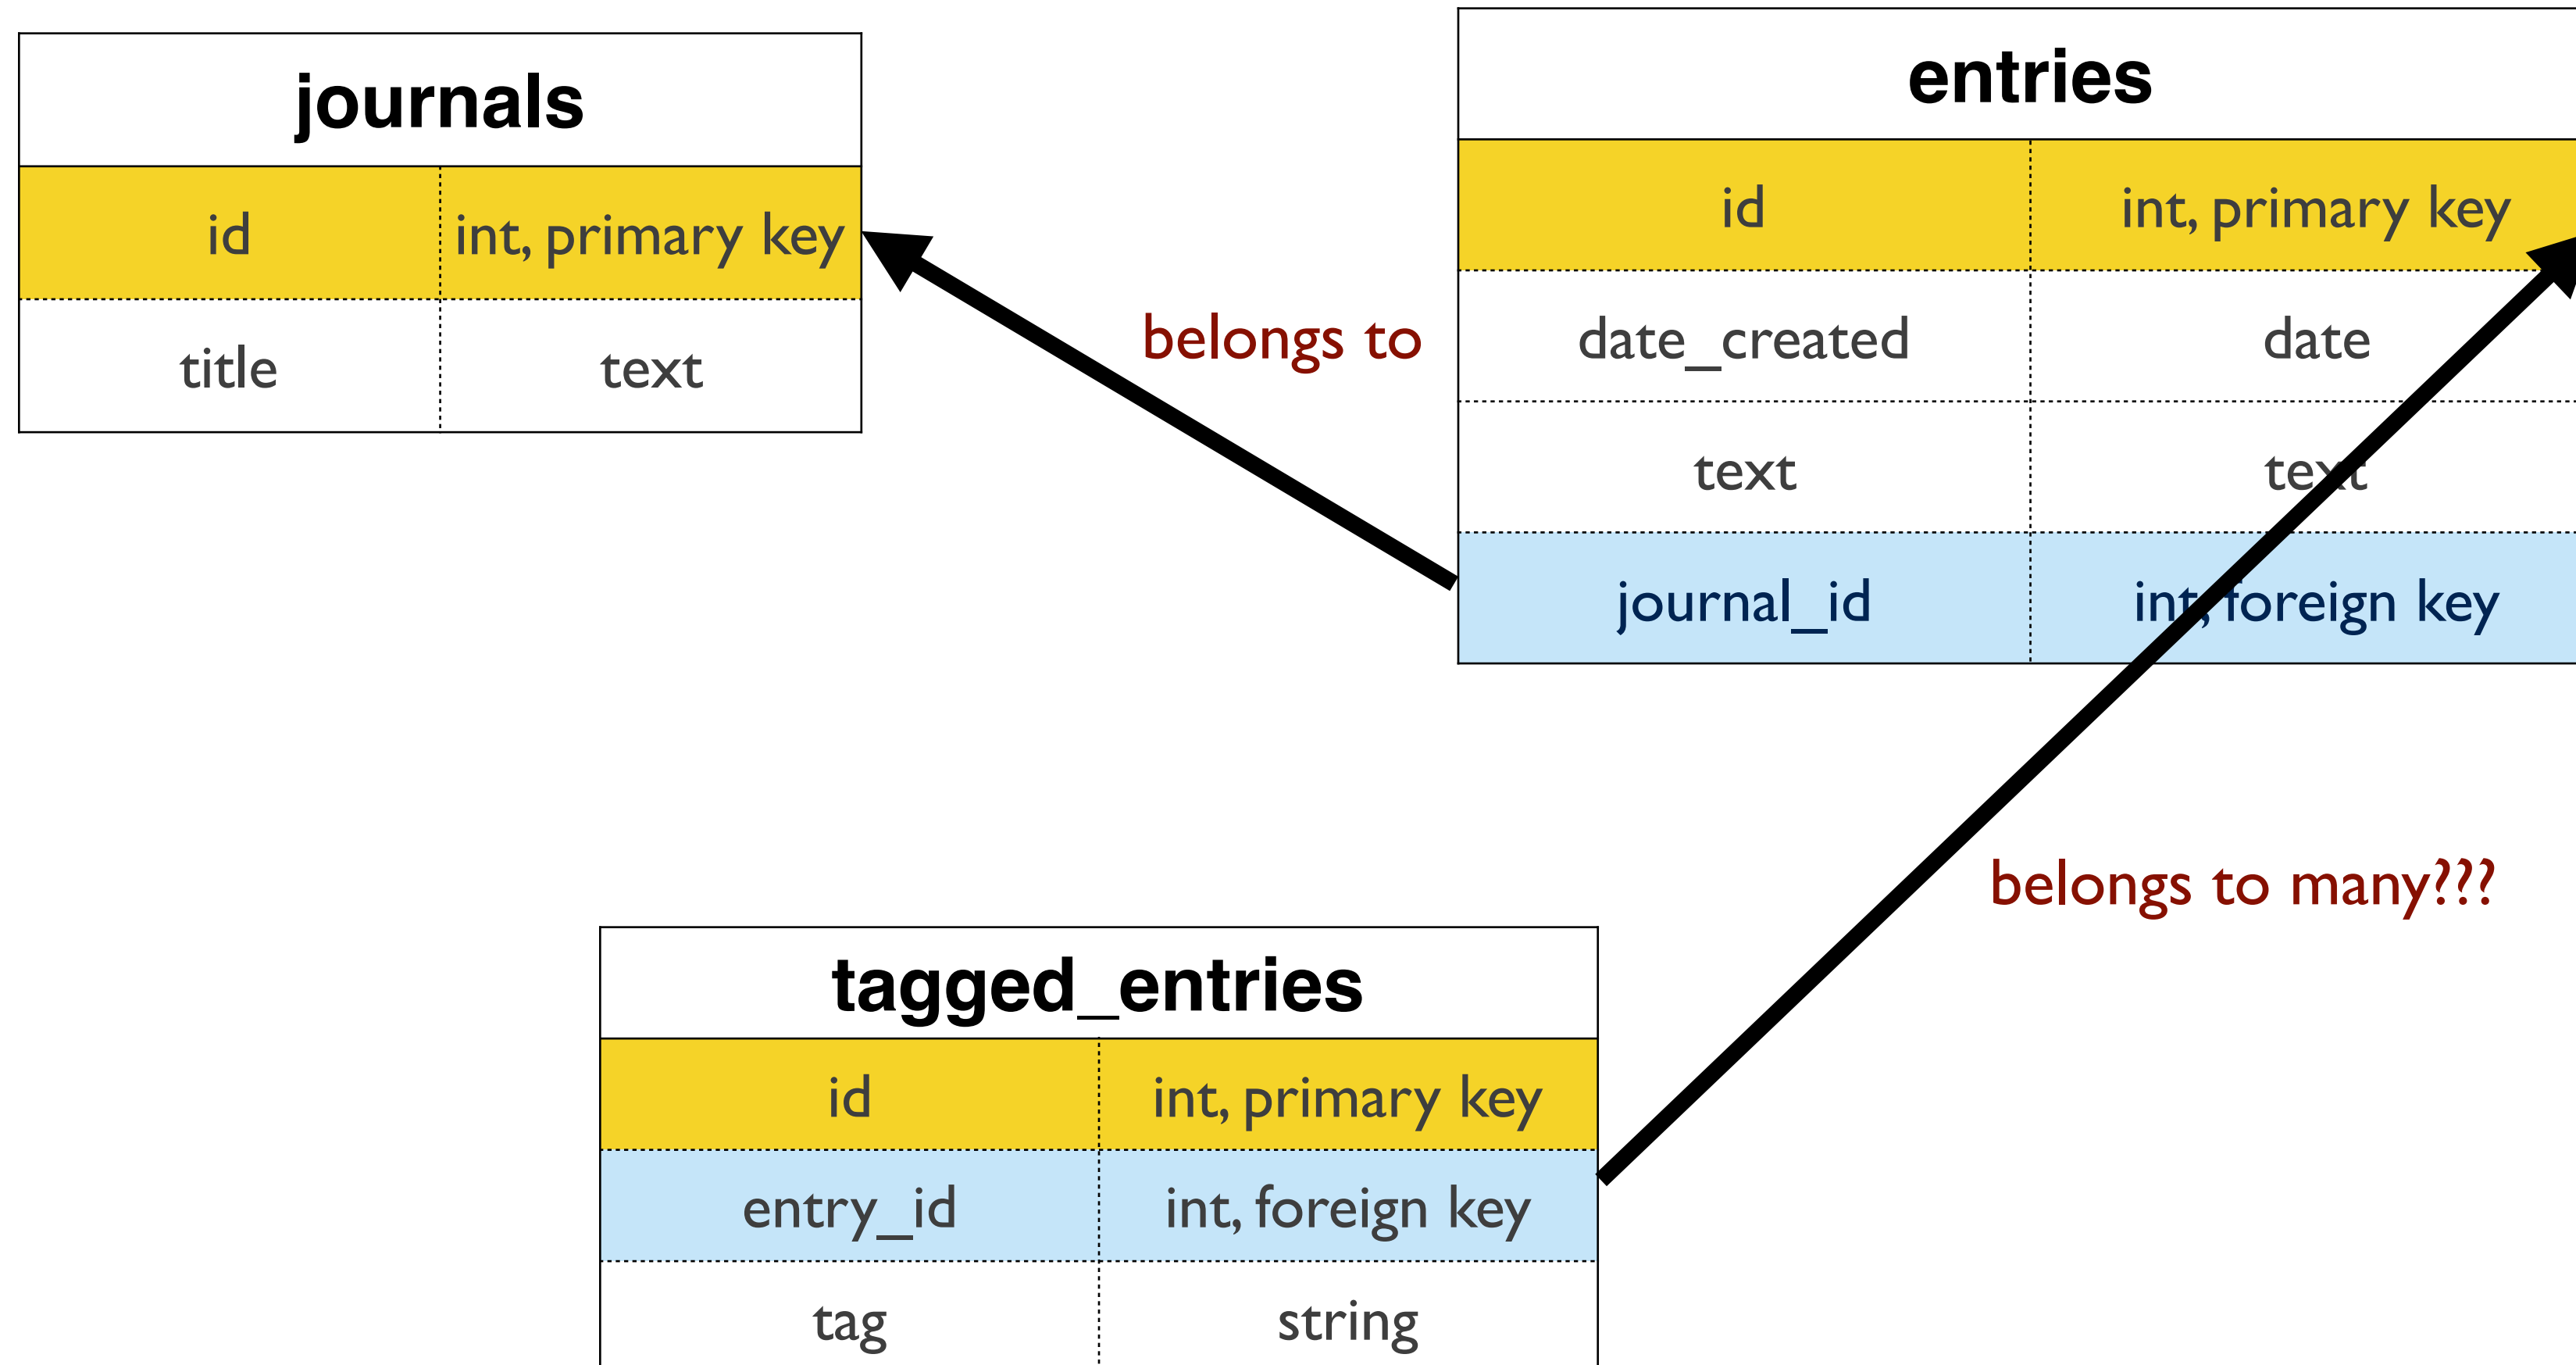
Example: A Journal

...But what about tags?!

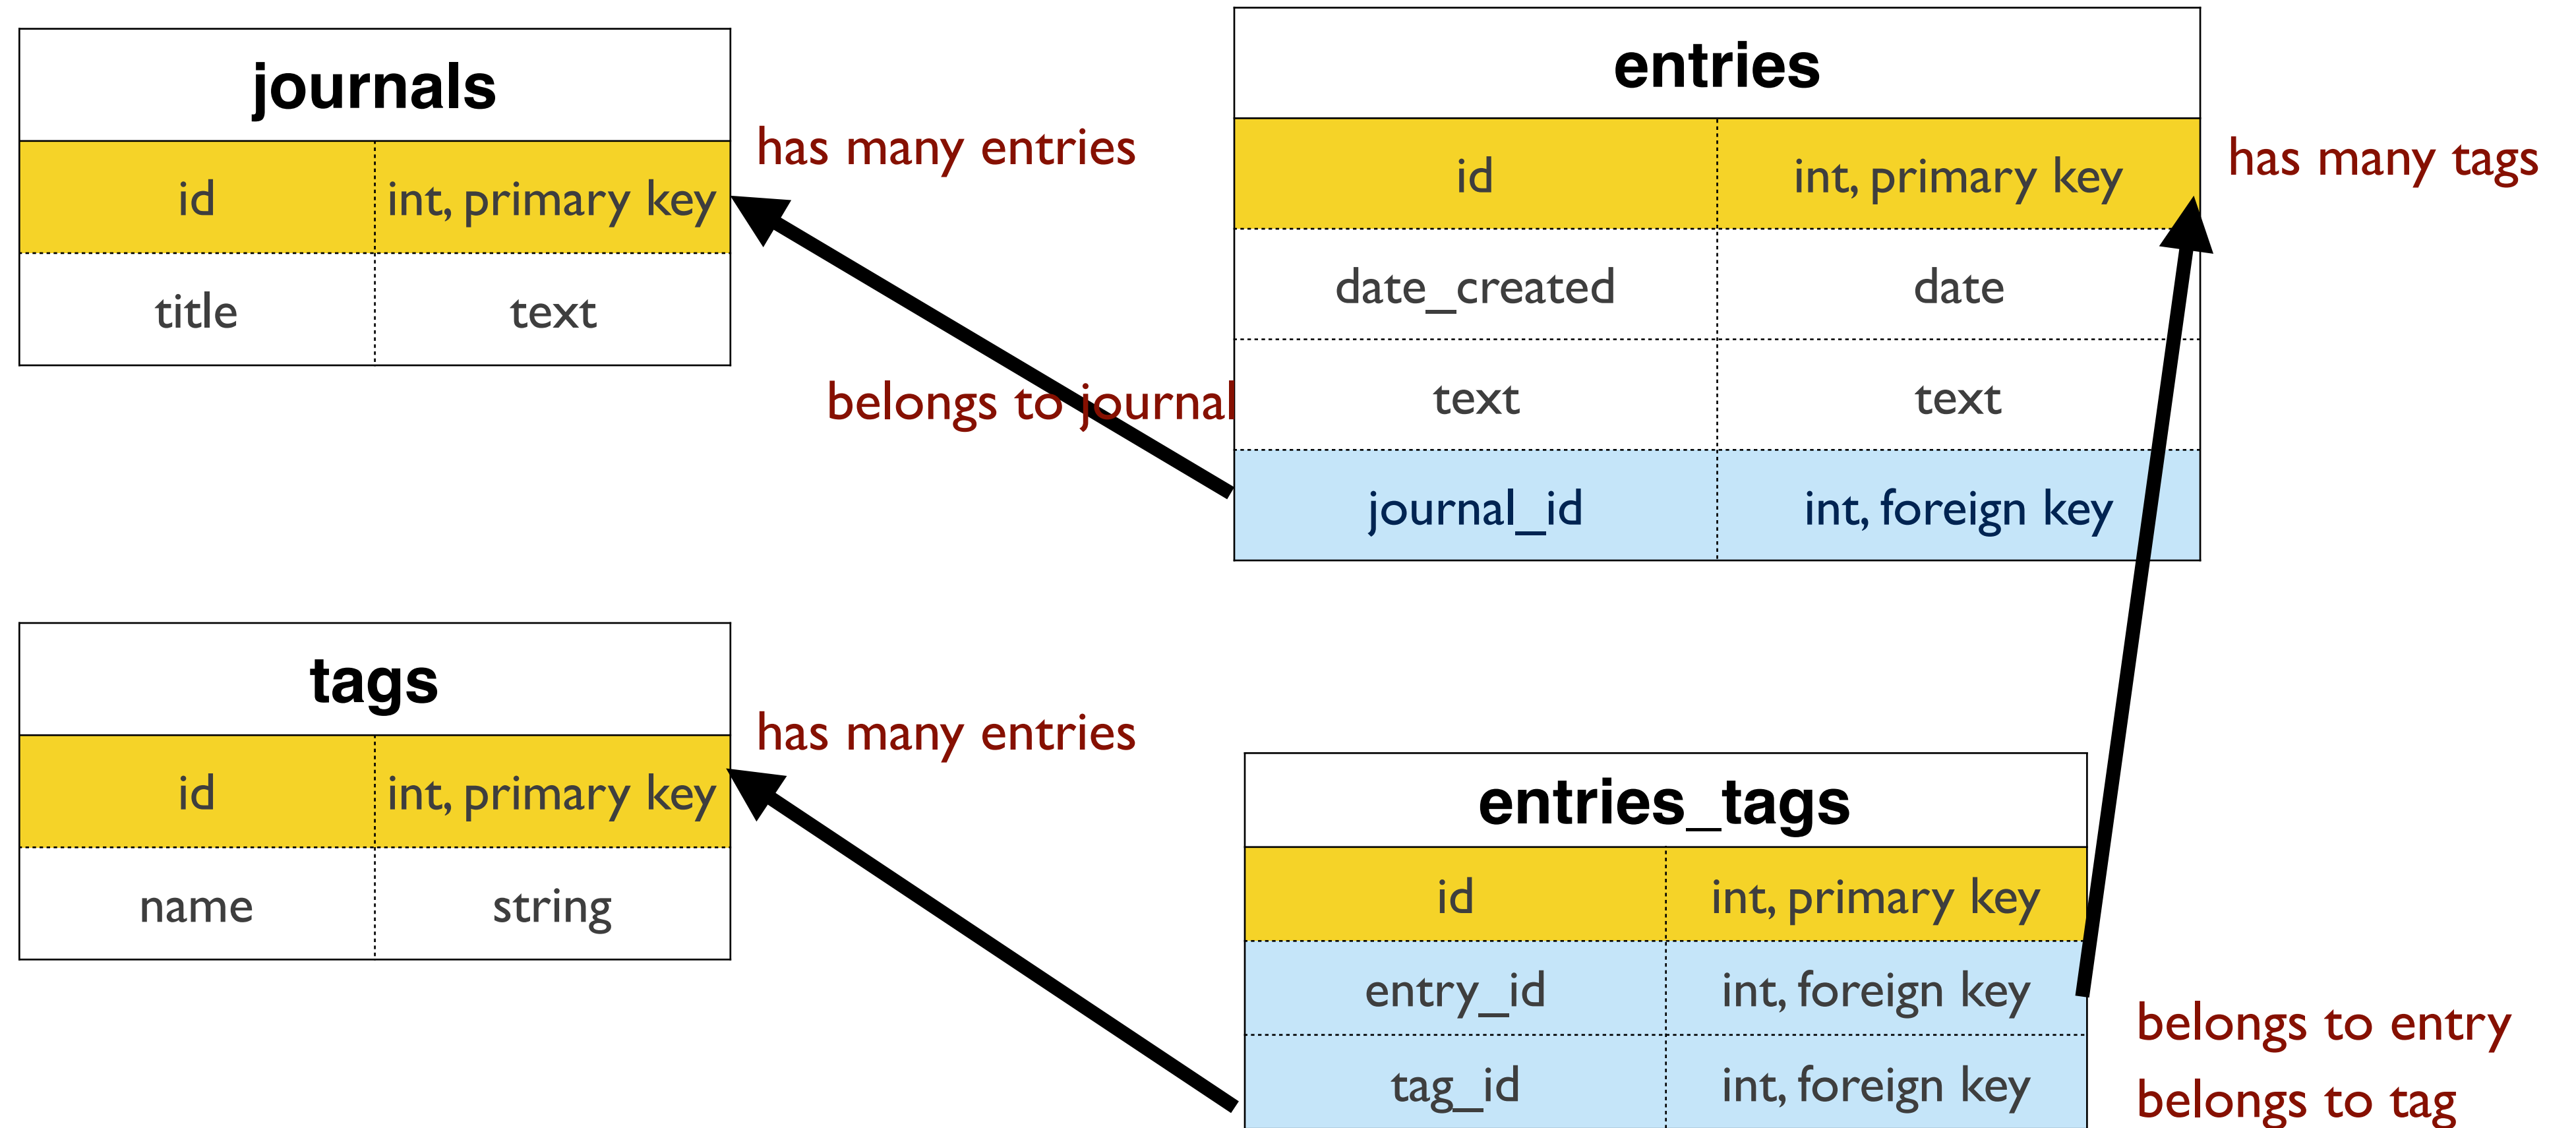


Example: A Journal

...What about now?



Example: A Journal Normalized!



Example: A Journal Normalized!

journals	
id	text
0	Fullstack journey
1	Work
2	Dank Memes
3	Despair is everywhere

tags	
id	text
0	#NoFilterNoMakeup
1	#YOLO
2	#selfie
3	#SadnessIsMyLife

entries			
id	date_created	text	journal_id
0	2016-04-01	I am happy	3
1	2016-04-02	I am very happy	2
2	2016-04-03	Great weekend	1
3	2016-04-03	Fun times	3

entries_tags		
id	tag_id	entry_id
0	2	1
1	3	1
2	1	1
3	3	2

Rule of Thumb

- **For a 1:M relationship:**

- One of the tables should have a foreign key column. Think about which thing 'belongs to' the other?
- E.g. dogs belong to owners. Each dog has a collar saying which owner it belongs to. The owner doesn't have a list of the dogs he/she owns.
- E.g. gloves belong to students. Every glove you have has your name in it. the glove table has the student foreign key in it.

- **For a M:M relationship:**

- You will need a *join table* which contains the foreign keys of the entities you are joining.

Normalized Databases

- Focus on optimal storage - often at odds with retrieval speed due to complex queries using complicated joins
- Work best when the application is write-intensive and write-load is more than read-load
 - Tables are usually smaller as data is divided vertically (fast reads on single tables)
 - Updates and Inserts are fast because there are no duplicates to update
 - Data is not duplicated so there is less of a need for process intensive group by or distinct queries
- Normalized tables mean join tables, which mean read operations on multiple tables suffer (indexing strategies don't work as well with joins)

Denormalized

- ⦿ **Works best when the application is read-intensive**
 - The data is present in the same table (no need for joins)
 - A single table with all required data allows for efficient index usage
- ⦿ **Data is duplicated which means that updates and inserts become complex and costly**

What Do I Do?!

- Real world applications will most likely have both read-loads and write-loads
- Utilize both approaches depending on the situation!
- Also, let your DBA handle most of this...



Design one!



- Twitter
- Gmail
- Facebook
- Instagram

- Wordpress
- Wikipedia
- AirBnB
- Google (search)



Steps for Developing your ERD

1. Identify Entities
2. Define Relationships
3. Draw Rough-Draft ERD
4. Fill in Cardinality/Modality (arrows with relationship type)
5. Define Primary Keys
6. Draw Key-Based ERD (labeling Primary and Foreign Keys)
7. Identify Attributes
8. Map Attributes
9. Draw fully attributed ERD